

OpenSecOps SOAR S3 Buckets

Standard Operating Procedure

Table of Contents

1 Introduction	3
2 Scope of SOP	3
Document Versions	3
3 Make your S3 Bucket Non-Public	4
3.1 CloudFormation	4
3.2 Terraform	4
4 Public S3 Buckets	5
5 Enabling S3 Bucket Server-Side Encryption	6
5.1 CloudFormation	6
5.2 Terraform	6
5.3 Buckets created manually	6
6 Other Considerations	7
6.1 Versioning	8
6.2 Life-Cycle Rules	8
6.3 Glacier Considerations	8
7 Security Hub Controls for S3 Buckets	9
7.1 CRITICAL	9
7.2 HIGH	9
7.3 MEDIUM	9
7.4 LOW	9

1 Introduction

The purpose of this document is to specify baseline settings for AWS S3 Buckets so that teams can provision S3 Buckets in a safe, conformant way.

2 Scope of SOP

This document covers AWS S3 Buckets and how they are to be set up in the CloudSecOps AWS organisation.

Document Versions

Version	Date	Changes	Author
1.0	2022-09-12	First version	Peter Bengtson
1.1	2025-04-07	Replaced "Delegat" with "OpenSecOps" throughout	Peter Bengtson

Please add new entries *above* already existing rows.

3 Make your S3 Bucket Non-Public

Always explicitly set your bucket to be non-public, unless it needs to be publicly accessible (for which see next section). Failing to make it inaccessible to the world is a **CRITICAL** security issue that immediately will be detected by our automated security monitoring. A security incident will be created which your team will have to fix as a show-stopping silver bullet, i.e. with immediate priority over all ongoing work in the team as a whole.

3.1 CloudFormation

```
ExampleBucket:
  Type: AWS::S3::Bucket
  Properties:
    PublicAccessBlockConfiguration:
      BlockPublicAcls: true
      BlockPublicPolicy: true
      IgnorePublicAcls: true
      RestrictPublicBuckets: true
```

3.2 Terraform

```
resource "aws_s3_bucket" "example" {
  bucket = "example"
}

resource "aws_s3_bucket_public_access_block" "example" {
  bucket = aws_s3_bucket.example.id

  block_public_acls    = true
  block_public_policy  = true
}
```

4 Public S3 Buckets

If you need an S3 bucket to be publicly readable and/or writable – which should be very rare – you must set one or both of the following tags on your S3 bucket:

```
soar:s3:request-publicly-readable  
soar:s3:request-publicly-writable
```

The tags are case-sensitive. The value of the tags should be empty (""). (NB: `publicly` is the correct spelling; "-ally" would be incorrect.)

When one or both of the above tags are present, the S3 bucket will be left open.

SOC will be notified of your request via a Jira incident. You will be contacted about the reasons for wanting to create a storage volume open to the world. Should SOC deny the request, you will also be contacted prior to the removal of the tag.

If the S3 bucket already has been closed by the Moonstone SOAR, adding the tags after the fact will *not* open the bucket again. You must add the tag(s), then open the bucket for read and/or write again either manually in the console, using the CLI, the API, or in code.

5 Enabling S3 Bucket Server-Side Encryption

All AWS S3 buckets created in Moonstone **MUST** have server-side encryption enabled. This applies to S3 buckets created manually, by using Terraform, or by using CloudFormation. It also applies to all S3 buckets created by or for third-party software such as Jenkins and Curity, as well as to buckets created by, and for, such things as CodePipeline.

Remediating this is very easy. Use the following tips to enable AES-256 encryption for existing buckets. Also, incorporate the code snippets in all your future infrastructure-as-code-created S3 buckets, so they are compliant right from the start, in all environments.

NB: Encryption can be enabled at any time. It doesn't matter whether the S3 bucket is empty or not. Encryption is completely transparent.

5.1 CloudFormation

```
ExampleBucket:
  Type: 'AWS::S3::Bucket'
  Properties:
    BucketEncryption:
      ServerSideEncryptionConfiguration:
        - ServerSideEncryptionByDefault:
            SSEAlgorithm: AES256
```

5.2 Terraform

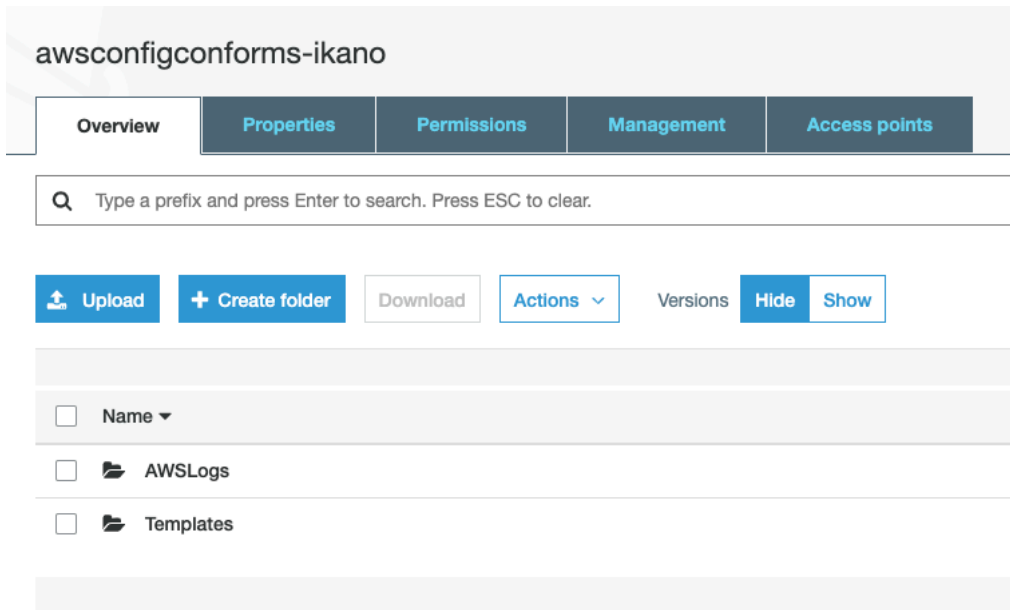
NB: The following example will encrypt the bucket and all its contents. If you have copied the S3 boilerplate example code from the Terraform documentation, you might have copied the example which encrypts all items in the bucket but not the bucket itself. That kind of setup will not satisfy the security control. The following, on the other hand, will:

```
resource "aws_s3_bucket" "mybucket" {
  bucket = "mybucket"

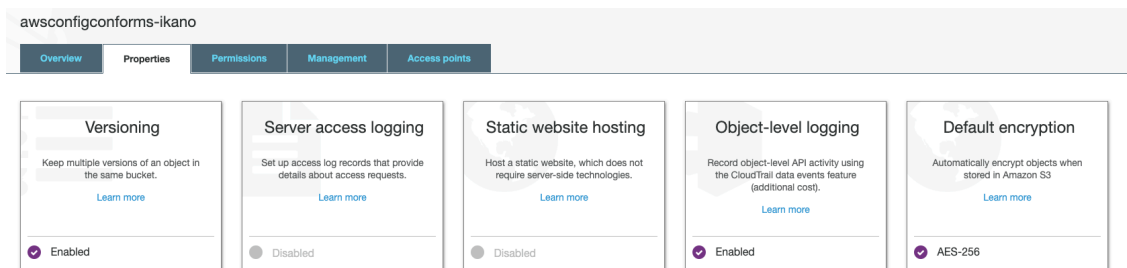
  server_side_encryption_configuration {
    rule {
      apply_server_side_encryption_by_default {
        sse_algorithm = "AES256"
      }
    }
  }
}
```

5.3 Buckets created manually

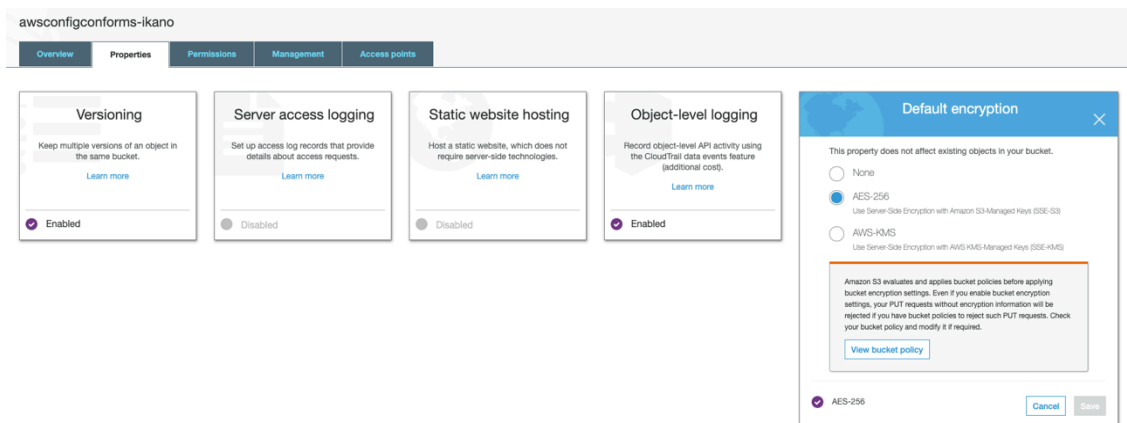
Use the console to navigate to the bucket, e.g.:



Click on Properties:



Then click on Default encryption and choose AES-256:



6 Other Considerations

6.1 Versioning

Consider enabling versioning for S3 buckets containing data which must not be lost or where changes to data must be trackable at all times. If you enable versioning, you should also enable one or more life-cycle rules for the data in the bucket.

6.2 Life-Cycle Rules

If your bucket contains logs or any continuously incoming high-volume data, consider enabling life-cycle rules to delete data after a suitable period of time, or storage costs might become an issue.

6.3 Glacier Considerations

If you write life-cycle rules to transition S3 bucket data to Glacier, make sure the items in your bucket aren't too small. There is a lower limit to what can be transferred to Glacier; also, retrieving small items from Glacier can become expensive. Therefore: do not store items smaller than 200 KB in Glacier.

7 Security Hub Controls for S3 Buckets

Below are the enabled controls in AWS Security Hub pertaining to S3 buckets. Your buckets must comply with all of them. Be proactive. It is easiest to make them compliant from the beginning, rather than when the automated security checks have created tickets for your team to fix misconfigurations which constitute security risks.

In the following list, all CIS.n.n required controls are satisfied by the enterprise infrastructural layer; you should never have to consider yourself with these.

However, the four S3.n rules (S3.2, S3.3, S3.4, and S3.6) are your responsibility at all times, no matter how you create your infrastructure.

7.1 CRITICAL

■ Critical	S3.2	S3 buckets should prohibit public read access
■ Critical	S3.3	S3 buckets should prohibit public write access
■ Critical	CIS.2.3	Ensure the S3 bucket used to store CloudTrail logs is not publicly accessible

7.2 HIGH

■ High	S3.6	S3 permissions granted to other AWS accounts in bucket policies should be restricted
--------	------	--

7.3 MEDIUM

■ Medium	S3.4	S3 buckets should have server-side encryption enabled
----------	------	---

7.4 LOW

■ Low	CIS.2.2	Ensure CloudTrail log file validation is enabled
■ Low	CIS.2.4	Ensure CloudTrail trails are integrated with CloudWatch Logs
■ Low	CIS.2.6	Ensure S3 bucket access logging is enabled on the CloudTrail S3 bucket
■ Low	CIS.3.8	Ensure a log metric filter and alarm exist for S3 bucket policy changes